



COMP2007/2907 - Algorithms

Course page: Blackboard and Ed

Lecturer: Joachim Gudmundsson
Level 4, Room 416, School of IT
joachim.gudmundsson@sydney.edu.au
Ph. 9351 4494

Tutors:	Xavi Holt (TA)	Yilun He
	Natalie Tridgell	Joe Godbehere
	Jessica McBroom	Patrick Eades
	Anton Jurisevic	Shumin Kong
	Gengxing Wang	Hisham Husein
	Mingshen Cai	



Course book:


J. Kleinberg and E. Tardos
Algorithm Design
Addison-Wesley

Outline:

12 lectures
5 assignments
10+1 quizzes
Exam

Tutorials:

13 tutorials



- › This unit provides an introduction to the design and analysis of algorithms. Its main aims are
 - (i) learn how to develop algorithmic solutions to computational problem
 - (ii) develop understanding of algorithm efficiency.

- › Assumes basic knowledge of discrete math
 - graphs
 - big O notation
 - proof techniquesand programming.



Assessment:

Quizzes 20% (average of best 8 out of 10)
Each assignment 6% (5 assignments - total 30%)
Exam 50% (minimum 40% required to pass)

Submissions:

Theory part - Blackboard (checked by Turnitin)
Implementation - Ed

Collaboration:

General ideas - Yes!
Formulation and writing - No!
Read [Academic Dishonesty and Plagiarism.](#)



- › There will be **5** homework assignments
- › The objective of these is to teach problem solving skills
- › Each assignment represents **6% of your final mark**. Late submissions will be penalized by 25% of the full marks per day.

For example, say you get 80% on your assignment:

If submitted on time = 4.8

Late but within 24 hours = $4.8 * 0.75 = 3.6$

Between 24 and 48 hours = $4.8 * 0.5 = 2.4$

Between 48 and 72 hours = $4.8 * 0.25 = 1.2$

More than 72 hours = $4.8 * 0 = 0$



- › The final will be 2.5 hours long. It will consist of 6 problems similar to those seen in the tutorials and assignments
- › The final will test your problem solving skills
- › There is a **40% exam barrier**
- › The final exam represents **50% of your final mark**
- › Our advice is that you work hard on the assignments throughout the semester. It's the best preparation for the final.



- › To get the most out of the tutorial, try to solve as many problems as you can *before* the tutorial. Your tutor is there to help you out if you get stuck, not to lecture.
- › We will post solutions to tutorials (see Ed).

- › **Lecture 1** [Mon 31 July]: Introduction
- › **Lecture 2** [Mon 7 Aug]: Graphs
- › **Lecture 3** [Mon 14 Aug]: Greedy algorithms
- › **Lecture 4** [Mon 21 Aug]: Divide & Conquer algorithms
- › **Lecture 5** [Mon 28 Aug]: Sweepline algorithms
- › **Lecture 6** [Mon 4 Sep]: Dynamic programming: basic techniques
- › **Lecture 7** [Mon 11 Sep]: Dynamic programming: interval scheduling and Bellman-Ford
- › **Lecture 8** [Mon 18 Sep]: Network flows I: Theory

Mon 25 Sep: University break

Mon 2 Oct: Labour Day

- › **Lecture 9** [Mon 9 Oct]: Network flows II: Applications
- › **Lecture 10** [Mon 16 Oct]: NP and intractability
- › **Lecture 11** [Mon 23 Oct]: Coping with hardness
- › **Lecture 12** [Mon 30 Oct]: Recap

Special Consideration (University policy)

- › If your performance on assessments is affected by illness or misadventure
 - › Follow proper bureaucratic procedures
 - Have professional practitioner sign special USyd form
 - Submit application for special consideration online, upload scans
 - Note you have only a quite short deadline for applying
 - http://sydney.edu.au/current_students/special_consideration/
 - › Also, notify coordinator by email *as soon as anything begins to go wrong*
 - › There is a similar process if you need special arrangements eg for religious observance, military service, representative sports
-

Academic dishonesty and plagiarism

- Please read the University policy on Academic Honesty carefully:
http://sydney.edu.au/elearning/student/EI/academic_honesty.shtml
- All cases of academic dishonesty and plagiarism will be investigated
- There is a new process and a centralized University system and database
- Three types of offenses:
 - **Plagiarism** – when you copy from another student, website or other source. This includes copying the whole assignment or only a part of it.
 - **Academic dishonesty** – when you make your work available to another student to copy (the whole assignment or a part of it). There are other examples of academic dishonesty.
 - **Misconduct** - when you engage another person to complete your assignment (or a part of it), for payment or not. This is a **very serious** matter and the Policy requires that your case is forwarded to the University Registrar for investigation.

- The penalties are **severe** and include:
 - 1) a permanent record of academic dishonesty, plagiarism and misconduct in the University database and on your student file
 - 2) mark deduction, ranging from 0 for the assignment to Fail for the course
 - 3) expulsion from the University and cancelling of your student visa
- **Do not confuse legitimate co-operation and cheating!** You can discuss the assignment with another student, this is a legitimate collaboration, but you cannot complete the assignment together – everyone must write their own code or report, unless the assignment is group work.
- When there is copying between students, note that **both students are penalised** – the student who copies and the student who makes his/her work available for copying

- We will use the similarity detection software TurnItIn and MOSS to compare your assignments with these of other students (current and previous) and the Internet
 - Turnitin is for text documents: http://www.turnitin.com/en_us/higher-education
 - MOSS is for programming code: <https://theory.stanford.edu/~aiken/moss/>
- These tools are **extremely good!**
 - e.g. MOSS cannot be fooled by changing the names of the variables or changing the order of the conditions in `if-else` statements
- Examples of plagiarism in programming code:
 - http://www.upenn.edu/academicintegrity/ai_computercode.html

- Plagiarism and any form of academic dishonesty will be dealt with, and the penalties are severe
- We use plagiarism detection systems such as MOSS that are extremely good. If you cheat, the chances you will be caught are very high.
- If someone asks you to see or copy your assignment, or to complete the assignment instead of them, just say: *I can't do this - we can both be thrown out of the University. I will not risk my future by doing this.*

Be smart and don't risk your future by engaging in plagiarism and academic dishonesty!



- › There are a wide range of support services available for students
 - › Please make contact, and get help
 - › You are not required to tell anyone else about this
-

Do you have a disability?

- › You may not think of yourself as having a ‘disability’ but the definition under the **Disability Discrimination Act** is broad and includes temporary or chronic medical conditions, physical or sensory disabilities, psychological conditions and learning disabilities.
 - › The types of disabilities we see include:
 - › anxiety, arthritis, asthma, asperger's disorder, ADHD, bipolar disorder, broken bones, cancer, cerebral palsy, chronic fatigue syndrome, crohn's disease, cystic fibrosis, depression, diabetes, dyslexia, epilepsy, hearing impairment, learning disability, mobility impairment, multiple sclerosis, post traumatic stress, schizophrenia , vision impairment, and much more.
 - › Students needing assistance must register with Disability Services –
 - › <http://sydney.edu.au/study/academic-support/disability-support.html>
-

- › Learning support
 - <http://sydney.edu.au/study/academic-support/learning-support.html>
 - › International students
 - <http://sydney.edu.au/study/academic-support/support-for-international-students.html>
 - › Aboriginal and Torres Strait Islanders
 - <http://sydney.edu.au/study/academic-support/aboriginal-and-torres-strait-islander-support.html>
 - › Student organization (can represent you in academic appeals etc)
 - <http://srcusyd.net.au/> or <http://www.supra.net.au/>
 - › Please make contact, and get help
 - › You are not required to tell anyone else about this
-



› Metacognition

- Pay attention to the learning outcomes in CUSP
- Self-check that you are achieving each one
- Think how each assessment task relates to these

› Time management

- Watch the due dates
- Start work early, submit early

› Networking and community-formation

- Make friends and discuss ideas with them
- Know your tutor and lecturer

› Enjoy the learning!

COMP2007/2907: Algorithms

Algorithms then, and now



- Algorithms can have huge impact
- For example -

A report to the White House from 2010 includes the following.

- Professor Martin Grotschel
 - A benchmark production planning model solved using linear programming would have taken 82 years to solve in 1988, using the computers and the linear programming algorithms of the day.
 - Fifteen years later, in 2003, this same model could be solved in roughly 1 minute, an improvement by a factor of roughly 43 million!
- [Extreme case, but even the average factor is very high.]
-

- In 2003 there were examples of problems that we can solve 43 million times faster than in 1988
 - This is because of better hardware and better algorithms

- In 1988
 - Intel 386 and 386SX
 - About 275,000 transistors
 - clock speeds of 16MHz, 20MHz, 25MHz, and 33MHz
 - MSDOS 4.0 and windows 2.0
 - VGA
 - In 2003
 - Pentium M
 - About 140 million transistors
 - Up to 2.2 GHz
 - AMD Athlon 64
 - Windows XP
-

- In a report to the White House from 2010 includes the following.
 - Professor Martin Grotschel:
 - A benchmark production planning model solved using linear programming would have taken 82 years to solve in 1988, using the computers and the linear programming algorithms of the day.
 - Fifteen years later, in 2003, this same model could be solved in roughly 1 minute, an improvement by a factor of roughly 43 millions

Observation:

- Hardware: 1,000 times improvement
 - Algorithms: 43,000 times improvement
-

- Efficient algorithms produce results within available resource limits
 - In practice
 - Low polynomial time algorithms behave well
 - Exponential running times are infeasible except for very small instances, or carefully designed algorithms
 - Performance depends on many obvious factors
 - Hardware
 - Software
 - Algorithm
 - Implementation of the algorithm
 - **This unit:** Algorithms
-

- Efficient algorithms “do the job” the way you want them to...
 - Do you need the exact solution?
 - Are you dealing with some special case and not with a general problem?
 - Is it ok if you miss the right solution sometimes?
-




- Complex, highly sophisticated algorithms can greatly improve performance

but...

- Reasonably good algorithmic solutions that avoid simple, or “lazy” mistakes, can have a much bigger impact!
-

List of topics

Greedy algorithms
Divide and conquer
Sweepline
Dynamic programming
Network flows
Mincut theorems
Approximation
Optimization problems




Introduction: Some Representative Problems

Four Representative Problems: Interval Scheduling

- › Input. Set of jobs with start times and finish times.
- › Goal. Find **maximum cardinality** subset of mutually compatible jobs.


↑
jobs don't overlap





Weighted Interval Scheduling


- › Input. Set of jobs with start times, finish times, and weights.
- › Goal. Find **maximum weight** subset of mutually compatible jobs.





Bipartite Matching

- › **Input.** Bipartite graph.
- › **Goal.** Find **maximum cardinality** matching.






Independent Set

- › Input. Graph.
- › Goal. Find **maximum cardinality** independent set.

↑
subset of nodes such that no two
joined by an edge




Four Representative Problems

- › Interval scheduling: $O(n \log n)$ time greedy algorithm.
- › Weighted interval scheduling: $O(n \log n)$ dynamic programming algorithm.
- › Bipartite matching: $O(n^3)$ maxflow based algorithm.
- › Independent set: NP-complete.



Algorithm Analysis & Data Structures



- › **Brute force.** For many non-trivial problems, there is a natural brute force search algorithm that checks every possible solution.
 - Typically takes 2^N time or worse for inputs of size N .
 - Unacceptable in practice.
- › **Desirable scaling property.** When the input size doubles, the algorithm should only slow down by some constant factor C .

There exists constants $c > 0$ and $d > 0$ such that on every input of size N , its running time is bounded by $c N^d$ steps.

- › **Definition:** An algorithm is **poly-time** if the above scaling property holds.

- › **Worst case running time.** Obtain bound on **largest possible** running time of algorithm on input of a given size N .
 - Generally captures efficiency in practice.
 - Draconian view, but hard to find effective alternative.

- › **Average case running time.** Obtain bound on running time of algorithm on **random** input as a function of input size N .
 - Hard (or impossible) to accurately model real instances by random distributions.
 - Algorithm tuned for a certain distribution may perform poorly on other inputs.

- › **Definition:** An algorithm is **efficient** if its running time is polynomial.
- › **Justification:** It really works in practice!
 - In practice, the poly-time algorithms that people develop almost always have low constants and low exponents.
 - Breaking through the exponential barrier of brute force typically exposes some crucial structure of the problem.
- › **Exceptions.**
 - Some poly-time algorithms do have high constants and/or exponents, and are useless in practice.
 - Some exponential-time (or worse) algorithms are widely used because the worst-case instances seem to be rare.

↑
simplex method
Unix grep



Table 2.1 The running times (rounded up) of different algorithms on inputs of increasing size, for a processor performing a million high-level instructions per second. In cases where the running time exceeds 10^{25} years, we simply record the algorithm as taking a very long time.

	n	$n \log_2 n$	n^2	n^3	1.5^n	2^n	$n!$
$n = 10$	< 1 sec	< 1 sec	< 1 sec	< 1 sec	< 1 sec	< 1 sec	4 sec
$n = 30$	< 1 sec	< 1 sec	< 1 sec	< 1 sec	< 1 sec	18 min	10^{25} years
$n = 50$	< 1 sec	< 1 sec	< 1 sec	< 1 sec	11 min	36 years	very long
$n = 100$	< 1 sec	< 1 sec	< 1 sec	1 sec	12,892 years	10^{17} years	very long
$n = 1,000$	< 1 sec	< 1 sec	1 sec	18 min	very long	very long	very long
$n = 10,000$	< 1 sec	< 1 sec	2 min	12 days	very long	very long	very long
$n = 100,000$	< 1 sec	2 sec	3 hours	32 years	very long	very long	very long
$n = 1,000,000$	1 sec	20 sec	12 days	31,710 years	very long	very long	very long

- › **Upper bounds.** $T(n)$ is $O(f(n))$ if there exist constants $c > 0$ and $n_0 \geq 0$ such that for all $n \geq n_0$ we have $T(n) \leq c \cdot f(n)$.
- › **Lower bounds.** $T(n)$ is $\Omega(f(n))$ if there exist constants $c > 0$ and $n_0 \geq 0$ such that for all $n \geq n_0$ we have $T(n) \geq c \cdot f(n)$.
- › **Tight bounds.** $T(n)$ is $\Theta(f(n))$ if $T(n)$ is both $O(f(n))$ and $\Omega(f(n))$.
- › **Ex:** $T(n) = 32n^2 + 17n + 32$.
 - $T(n)$ is $O(n^2)$, $O(n^3)$, $\Omega(n^2)$, $\Omega(n)$, and $\Theta(n^2)$.
 - $T(n)$ is not $O(n)$, $\Omega(n^3)$, $\Theta(n)$, or $\Theta(n^3)$.



- › Slight abuse of notation. $T(n) = O(f(n))$.
 - Asymmetric:
 - $f(n) = 5n^3$; $g(n) = 3n^3$
 - $f(n) = O(n^3) = g(n)$
 - but $f(n) \neq g(n)$.
 - Better notation: $T(n) \in O(f(n))$.

- › **Meaningless statement.** Any comparison-based sorting algorithm requires at least $O(n \log n)$ comparisons.
 - Statement doesn't "type-check."
 - Use Ω for lower bounds.



› Transitivity

- If $f = O(g)$ and $g = O(h)$ then $f = O(h)$.
- If $f = \Omega(g)$ and $g = \Omega(h)$ then $f = \Omega(h)$.
- If $f = \Theta(g)$ and $g = \Theta(h)$ then $f = \Theta(h)$.

› Additivity

- If $f = O(h)$ and $g = O(h)$ then $f + g = O(h)$.
- If $f = \Omega(h)$ and $g = \Omega(h)$ then $f + g = \Omega(h)$.
- If $f = \Theta(h)$ and $g = O(h)$ then $f + g = \Theta(h)$.



Asymptotic Bounds for Some Common Functions

- › **Polynomials.** $a_0 + a_1n + \dots + a_dn^d$ is $\Theta(n^d)$ if $a_d > 0$.
- › **Polynomial time.** Running time is $O(n^d)$ for some constant d independent of the input size n .

- › **Logarithms.** $O(\log_a n) \overset{\uparrow}{=} O(\log_b n)$ for any constants $a, b > 0$.
can avoid specifying the base

- › **Logarithms.** For every $x > 0$, $\log n \overset{\uparrow}{=} O(n^x)$.
log grows slower than every polynomial

- › **Exponentials.** For every $r > 1$ and every $d > 0$, $n^d \overset{\uparrow}{=} O(r^n)$.


every exponential grows faster than every polynomial



- › **Linear time.** Running time is at most a constant factor times the size of the input.
- › Computing the maximum. Compute maximum of n numbers a_1, \dots, a_n .

```
max ← a1
for i = 2 to n
{
    if (ai > max)
        max ← ai
}
```

- › **Merge.** Combine two sorted lists $A = a_1, a_2, \dots, a_n$ with $B = b_1, b_2, \dots, b_n$ into one sorted list.



```
i = 1, j = 1
while (both lists are nonempty) {
    if (a_i ≤ b_j) then append a_i to output list and increment i
    else append b_j to output list and increment j
}
append remainder of nonempty list to output list
```



- › $O(n \log n)$ time. Arises in divide-and-conquer algorithms.
- › Sorting. Mergesort and heapsort are sorting algorithms that perform $O(n \log n)$ comparisons.
- › Largest empty interval. Given n time-stamps x_1, \dots, x_n on which copies of a file arrive at a server, what is largest interval of time when no copies of the file arrive?
- › $O(n \log n)$ solution. Sort the time-stamps. Scan the sorted list in order, identifying the maximum gap between successive time-stamps.



Quadratic Time: $O(n^2)$

- › **Quadratic time.** Enumerate all pairs of elements.
- › Closest pair of points. Given a list of n points in the plane $(x_1, y_1), \dots, (x_n, y_n)$, find the pair that is closest.
- › $O(n^2)$ solution. Try all pairs of points.

```
min ←  $(x_1 - x_2)^2 + (y_1 - y_2)^2$  ← don't need to  
for i = 1 to n {  
    for j = i+1 to n {  
        d ←  $(x_i - x_j)^2 + (y_i - y_j)^2$   
        if (d < min)  
            min ← d ← see chapter 5  
    }  
}
```

- › **Cubic time.** Enumerate all triples of elements.
- › **Set disjointness.** Given n sets S_1, \dots, S_n each of which is a subset of $1, 2, \dots, n$, is there some pair of these which are disjoint?
- › $O(n^3)$ solution. For each pairs of sets, determine if they are disjoint.

```
foreach set  $S_i$  {  
    foreach other set  $S_j$  {  
        foreach element  $p$  of  $S_i$  {  
            determine whether  $p$  also belongs to  $S_j$   
        }  
        if (no element of  $S_i$  belongs to  $S_j$ )  
            report that  $S_i$  and  $S_j$  are disjoint  
    }  
}
```

Polynomial Time: $O(n^k)$ Time

- › Independent set of size k . Given a graph, are there k nodes such that no two are joined by an edge?
↖
 k is a constant
- › $O(n^k)$ solution. Enumerate all subsets of k nodes.

```
foreach subset S of k nodes {  
    check whether S is an independent set  
    if (S is an independent set)  
        report S is an independent set  
    }  
}
```

- Check whether S is an independent set = $O(k^2)$.

- Number of k element subsets = $\binom{n}{k} = \frac{n(n-1)(n-2)\cdots(n-k+1)}{k(k-1)(k-2)\cdots(2)(1)} \leq \frac{n^k}{k!}$

- $O(k^2 n^k / k!) = O(n^k)$.

poly-time for fixed k
but not practical
↙



- › Independent set. Given a graph, what is maximum size of an independent set?
- › $O(n^2 2^n)$ solution. Enumerate all subsets.

```
S* ←  $\phi$ 
foreach subset S of nodes {
    check whether S is an independent set
    if (S is largest independent set seen so far)
        update S* ← S
    }
}
```




Summary: Algorithm analysis

- › You must learn the asymptotic order of growth. It is fundamental when measuring the performance of an algorithm.
 - O -notation
 - Ω -notation
 - Θ -notation
- › Transitivity and additivity



Basic dynamic data structures

Assumed knowledge:

- Linked lists
- Queues
- Stacks
- Balanced binary trees




Why data structures?

- › Programs manipulate data
 - › Data should be organized so manipulations will be efficient
 - Search (e.g. Finding a word/file/web page)
 - › Good programs are powered by good data structures
 - › Naïve choices are often much less efficient than clever choices
 - › Data structures are existing tools that can help you
 - guide your design, and
 - save you time (avoid re-inventing the wheel)
-

The Queue data structure

- › The **Queue** data structure stores arbitrary objects
 - › Insertions and deletions follow the first-in first-out (FIFO) scheme
 - › Insertions are at the rear of the queue and removals are at the front of the queue
 - › Main queue operations:
 - **enqueue**(object): inserts an element at the end of the queue
 - object **dequeue**(): removes and returns the element at the front of the queue
 - › Auxiliary queue operations:
 - object **front**(): returns the element at the front without removing it
 - integer **size**(): returns the number of elements stored
 - boolean **isEmpty**(): indicates whether no elements are stored
-

- › The **Stack** data structure stores arbitrary objects
 - › Insertions and deletions follow the last-in first-out (LIFO) scheme
 - › Think of a spring-loaded plate dispenser
 - › Main stack operations:
 - **push**(object): inserts an element
 - object **pop**(): removes and returns the last inserted element
 - › Auxiliary stack operations:
 - object **top**(): returns the last inserted element without removing it
 - integer **size**(): returns the number of elements stored
 - boolean **isEmpty**(): indicates whether no elements are stored
-






Tree Terminology

- › Root: node without parent (A)
- › Internal node: node with at least one child (A, B, C, F)
- › External node (a.k.a. leaf): node without children (E, I, J, K, G, H, D)
- › Ancestors of a node: parent, grandparent, grand-grandparent, etc.
- › Depth of a node: number of ancestors
- › Height of a tree: maximum depth of any node (3)
- › Descendant of a node: child, grandchild, grand-grandchild, etc.

- › Subtree: tree consisting of a node and its descendants






› Notation

n number of nodes

e number of external nodes


i number of internal nodes

h height



Properties (binary trees):

- $e = i + 1$
- $n = 2e - 1$
- $h \leq i$
- $h \leq (n - 1)/2$
- $e \leq 2^h$
- $h \geq \log_2 e$
- $h \geq \log_2 (n + 1) - 1$





Running Times for AVL Trees

- › find is $O(\log n)$
 - height of tree is $O(\log n)$, no restructures needed

 - › insert is $O(\log n)$
 - initial find is $O(\log n)$
 - Restructuring up the tree, maintaining heights is $O(\log n)$

 - › remove is $O(\log n)$
 - initial find is $O(\log n)$
 - Restructuring up the tree, maintaining heights is $O(\log n)$
-



› Queues

- Enqueue, dequeue, first and size operations in $O(1)$ time.

› Stacks

- Push, pop, top and size operations in $O(1)$ time

› Balanced binary trees (e.g. AVL trees)

- Insert, delete and find operations in $O(\log n)$ time
-